Noisy Self-Training with Data Augmentations for Offensive and Hate Speech Detection Tasks

João A. Leite¹ and Carolina Scarton¹ and Diego F. Silva²

¹Department of Computer Science, The University of Sheffield, Sheffield (UK)

²Instituto de Ciências Matemáticas e de Computação, Universidade de São Paulo, São Carlos (Brazil)

{ jaleitel, c.scarton}@sheffield.ac.uk, diegofsilva@usp.br

Abstract

Online social media is rife with offensive and hateful comments, prompting the need for their automatic detection given the sheer amount of posts created every second. Creating highquality human-labelled datasets for this task is difficult and costly, especially because nonoffensive posts are significantly more frequent than offensive ones. However, unlabelled data is abundant, easier, and cheaper to obtain. In this scenario, self-training methods, using weakly-labelled examples to increase the amount of training data, can be employed. Recent "noisy" self-training approaches incorporate data augmentation techniques to ensure prediction consistency and increase robustness against noisy data and adversarial attacks. In this paper, we experiment with default and noisy self-training using three different textual data augmentation techniques across five different pre-trained BERT architectures varying in size. We evaluate our experiments on two offensive/hate-speech datasets and demonstrate that (i) self-training consistently improves performance regardless of model size, resulting in up to +1.5% F1-macro on both datasets, and (ii) noisy self-training with textual data augmentations, despite being successfully applied in similar settings, decreases performance on offensive and hate-speech domains when compared to the default method, even with state-ofthe-art augmentations such as backtranslation.

1 Introduction

Online social media platforms are widely used by modern society for many productive purposes. However, they are also known for intensifying offensive and hateful comments, attributed in part to factors such as user anonymity (Mondal et al., 2017). Manual identification of hate speech is impractical at scale due to the massive number of posts generated every second and the potential harm to the mental health of moderators. Therefore, there is a need for automatic approaches to detect offensive and hateful speech.

In recent years, research on this topic has increased, resulting in new models and datasets published in various languages and sources (Fortuna and Nunes, 2018). A common characteristic among available datasets is label skewness towards the negative class (non-offensive/hateful), which is usually more frequent than the positive class (offensive/hateful). Apart from traditional ways of dealing with imbalanced classes (e.g. under or oversampling or applying class weighting), semi-supervised techniques such as self-training can be used to extend the training set with unseen examples that introduce new learning signals without the costly burden of manual data labeling.

Self-training is a technique that involves iteratively training models using both labelled and unlabelled data. The process begins by training a model using human-labelled data only, which is then used to infer labels for a set of unlabelled data, creating a weakly-labelled dataset. The weakly-labelled dataset and the human-labelled dataset are then aggregated and used to retrain the model. This iterative process is repeated for a fixed number of steps or until no performance improvement is observed. Self-training can be particularly useful when labelled data is scarce or expensive to obtain, and was successfully applied in a variety of domains such as computer vision (Schiappa et al., 2022), audio and speech processing (Liu et al., 2022), and natural language processing (He et al., 2019).

Several variants of self-training have been proposed over the years (Amini et al., 2022). One common approach is to use a teacher-student framework, in which the "student" model learns from the output generated by the "teacher" model (Blum and Mitchell, 1998; Xie et al., 2020b; Chen et al., 2021; Karamanolakis et al., 2021). Additionally, a confi-

dence threshold filter may be applied to remove examples that are too ambiguous or non-informative. This process is summarised in Figure 1.

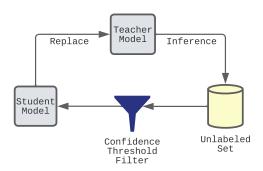


Figure 1: Teacher-student self-training loop

Recent research on self-training has reported further improvements in performance by introducing perturbations directly into the raw input or to its latent representation, improving generalisation and convergence (Rasmus et al., 2015; Laine and Aila, 2017; Miyato et al., 2018; He et al., 2019; Xie et al., 2020a). These perturbations are often introduced in the form of data augmentations, which are widely applied in Computer Vision tasks but are less commonly explored in Natural Language Processing tasks, especially in the context of selftraining. These "noisy self-training" methods can be particularly useful in settings where the input data is noisy or subject to a high degree of variation, improving prediction consistency and adversarial robustness (Carmon et al., 2019; Alayrac et al., 2019; Najafi et al., 2019).

Bayer et al. (2022) argue that data augmentation depends on the underlying classification task, thus it cannot be effectively applied in all circumstances. Previous work focusing solely on data augmentation methods, not coupled with self-training, has shown mixed results for the domain of offensive/hate speech classification (Section 2.1). This indicates that there may not be a best method, while some may even negatively impact performance.

An open question is whether noisy self-training with text data augmentations can contribute to text classification tasks using state-of-the-art transfer-learning BERT models that have been shown to be invariant to various data transformations (Long-pre et al., 2020). The task of offensive/abusive speech detection poses a difficult challenge for generating high-quality semantic invariant augmented examples, since it is a domain that is intrinsically associated with specific keywords that, if modified,

can completely change the semantics of the text. In this paper, we innovate by providing an extensive experimentation setup using three different data augmentation techniques - backtranslation, random word swap, and random synonym substitution - in a self-training framework, with five different pretrained BERT architectures varying in size, on two different datasets.

We demonstrate that self-training, either with or without data noising, outperforms default fine-tuning regardless of model size, on both datasets. However, when comparing self-training without data noising vs 'noisy' self-training, we find that data augmentations decrease performance, despite the literature reporting the superiority of noisy self-training in other domains. We further investigate how the augmentation methods fail to create label-invariant examples for the offensive/hate speech domain. Finally, we discuss future research ideas to address the limitations found in this work.

2 Related Work

2.1 Data Augmentation

Bayer et al. (2022) present a survey on data augmentation methods for NLP applications, reporting performance gains on various tasks.

In the domain of offensive/hate speech classification, Ibrahim et al. (2018) experiment with three different text augmentation techniques to expand and balance their Wikipedia dataset by augmenting negative (non-offensive) examples. From a binary view of the dataset, more than 85% of their examples are labelled as non-offensive, and from a multi-label view of the dataset, three of the six offensive classes are represented by less than 7% of the dataset. They report F1-score increases of +1.4% with unique words augmentation, +2.9% with unique words and random mask, and +3.6% with unique words, random mask, and synonym replacement.

Mosolova et al. (2018) use a custom synonym replacement augmentation method to experiment with a 'toxic' dataset with 6 classes from a Kaggle competition¹. They experiment with character and word embeddings with a CNN architecture, and report a +3.7% and +5.1% ROC-AUC increase when applying their augmentation method with character embeddings on the public and private

https://www.kaggle.com/c/
jigsaw-toxic-comment-classification-challenge

scores², respectively. However, when coupled with word embeddings, they find that their augmentations result in a decrease of -0.09% and -0.21% ROC-AUC scores on the public and private scores, respectively.

Rizos et al. (2019) propose three text-based data augmentation techniques to address the class imbalance in datasets, and apply them on three English hate speech datasets named HON (Davidson et al., 2017), RSN-1 (Waseem and Hovy, 2016) and RSN-2 (Waseem, 2016). Their augmentation methods include (i) synonym replacement based on word embedding, (ii) warping of the token words along the padded sequence, and (iii) class-conditional RNN language generation. They compare the three methods on different architectures combining word embeddings, CNNs, GRUs, and LSTMs, and they report an average across four different architecture configurations of -6.3% F1-Macro using (i), +5% F1-Macro using (ii), and -4% F1-Macro using (iii).

Marivate and Sefara (2020) experiment with four different data augmentation techniques: Word-Net synonym substitution, backtranslation between German and English, word embedding substitution according to cosine similarity, and mixup (Zhang et al., 2018). Authors experiment with three datasets from different domains: Sentiment 140 (Go et al., 2009), AG News (Zhang et al., 2015) and a Hate Speech dataset (Davidson et al., 2017). They observe performance increases on both Sentiment 140 and AG News across different augmentation methods, up to +0.4% and +0.5% accuracy score on AG News and Sentiment 140, respectively. However, they report performance decreases with all methods on the Hate Speech dataset, with decreases of 0.0% with mixup, -0.3% with embedding similarity, -0.8% with synonym substitution, and -2.3% with backtranslation.

2.2 Self-Training

Xie et al. (2020b) present a method called *noisy student*, which achieves state-of-the-art results on the ImageNet dataset (Deng et al., 2009) by performing self-training with a teacher-student approach, using student models that are equal or larger-sized than the teacher models, and adding noise both to the input data through random image augmentations and to the model via dropout.

He et al. (2019) apply a similar idea using textual

data augmentation methods such as backtranslation (Edunov et al., 2018) and token modifications to a self-training LSTM architecture for the tasks of machine translation and text summarization. They find that both model noise, in the form of dropout, and data noise, in the form of data augmentations, are crucial to their observed increase in performance on both tasks.

Xie et al. (2020a) use six text classification and two image classification benchmark datasets to experiment with different types of noise-inducing techniques for self-training. They argue that state-of-the-art augmentations like backtranslation for text classification and RandAugment (Cubuk et al., 2020) for image classification, outperform simple noise inducing techniques, such as additive Gaussian noise.

The use of noisy self-training approaches in the domain of offensive/hate speech classification is still limited, but default 'non-noisy' self-training has been successfully applied in some recent works. Alsafari and Sadaoui (2021) collect unlabelled Arabic tweets and perform semi-supervised classification with self-training for the domain of Offensive and Hate Speech detection using multiple text representations such as N-grams, Word2Vec, AraBert and Distilbert, and multiple model architectures such as SVM, CNN and BiLSTM. They report up to 7% performance increase in low resource settings where only a few labelled examples are available.

Leonardelli et al. (2020) apply self-training in their submission to the HaSpeeDe shared task on Italian hate speech detection (task A). They fine-tune an AlBERTo model with the human-labelled dataset provided by the task organisers and extend it with a weakly-labelled dataset using self-training. Additionally, they oversample the human-labelled set in an attempt to make the model more robust to inconsistencies in the weakly-labelled set. Their submission achieve an F1-macro score of 75.3% on tweets, placing 11th out of 29 teams, and 70.2% on news headlines, placing 5th out of 29 teams.

Pham-Hong and Chokshi (2020) report experiments with the noisy student method from Xie et al. (2020b) in the OffensEval 2020 shared task, achieving 2nd place at subtask B (Automatic categorization of offense types). In their setup, although dropout is applied to a BERT-large model, no noise is injected into the data, which is a crucial component of the noisy student method. Because of

²Public scores are computed over a smaller portion of the test set. At the end of the competition, private scores are computed with the remainder of the test set.

this, we argue that this work is actually applying a default self-training method instead of a noisy self-training method. Also, OffensEval 2020's training data does not contain human-labelled data³, thus both their weakly-labelled dataset and ground-truth dataset consist of inferred examples.

Richardson et al. (2022) detect hate speech on Twitter in the context of the Covid-19 pandemic. They employ a simple approach, utilizing a bag-of-words representation combined with an SVM classifier. Authors demonstrate that by employing self-training with only 20% of the training data, they manage to improve accuracy by +1.55% compared to default training using 80% of the training data.

To the best of our knowledge, Santos et al. (2022) is the only previous work in which a **noisy** self-training approach was attempted on an offensive/hate speech classification task. They propose an ensemble of two semi-supervised models to create FIGHT, a Portuguese hate speech corpus. Authors combine GANs, a BERT-based model, and a label propagation model, achieving 66.4% F1-score. They attempt to increase performance using backtranslation as data augmentation, but ultimately observe no performance gains, thus their best model is obtained with default self-training, not with noisy self-training.

3 Materials and Methods

This section presents the description of the datasets, data augmentation methods and self-training architectures used throughout our experiments. Our code is available at GitHub⁴.

3.1 Data Description

We use two English binary offensive/hate speech detection datasets in our experiments. Table 1 presents their target class distributions.

Offensive Language Identification Dataset (OLID) (Zampieri et al., 2019) contains a collection of annotated tweets following three levels: Offensive Language Detection, Categorization of Offensive Language, and Offensive Language Target Identification. This work only uses the first level - Offensive Language Detection. The dataset was

OLID			
	Train	Dev	Test
Not-Offensive	8,840	0	620
Offensive	4,400	0	240
Con	nvAbuse	:	
	Train	Dev	Test
Not-Offensive	2,163	719	725
Offensive	338	112	128

Table 1: Target class distribution for OLID and ConvAbuse.

normalised by replacing URLs and user mentions with placeholders. The best model in (Zampieri et al., 2019) achieves 80% macro-F1 using convolutional neural networks, with 70% and 90% of F1-Score for the positive and negative classes, respectively.

ConvAbuse (Cercas Curry et al., 2021) is a dataset on abusive language towards three conversational AI systems: an open-domain social bot, a rule-based chatbot, and a task-based system. Authors find that the distribution of abuse towards conversational systems differs from other commonly used datasets, with more than 50% of the instances containing sexism or sexual harassment. To normalise the data, web addresses were replaced with a placeholder. Authors provide standard train, development, and test sets and achieve up to 88.92% macro-F1 using a fine-tuned BERT model. In our experiments, we concatenate the interactions between the user and the chatbot into a single text document divided by new line separators, and we use majority voting between the annotations to consolidate the binary abusive vs. non-abusive label.

Unlabelled data We collected 365,456 tweets in English with the Twitter API using an unbiased query rule: random tweets mentioning stop-words like "in", "on", "a", "is", "not", "or" and so on. We also preprocess the data by removing user mentions, urls, punctuations, extra whitespace and accents.

3.2 Self-Training Architecture

Our noisy self-training system is similar to that introduced by Xie et al. (2020b) and Xie et al. (2020a), and works as follows:

 A teacher model is trained to minimise the cross-entropy loss on the human-labelled training set exclusively.

³In OffensEval 2020, the labels in the training data are the average confidence score and confidence standard deviation aggregated from an ensemble of models.

⁴https://github.com/JAugusto97/ Offense-Self-Training

- 2. The teacher model infers weak labels from the unlabelled dataset.
 - A confidence threshold filter is applied, and examples that fall below this threshold are removed.
 - Apply downsampling on the inferred examples, ending up with a perfectly balanced weakly-labelled dataset.
- 3. All the examples selected from the previous step are augmented once with one of the data augmentation methods, doubling the amount of weakly-labelled examples. The labels obtained with the 'clean/without noise' text in step 2 are replicated for the augmented texts.
- 4. An equal-sized student model minimises the combined cross-entropy loss on human-labelled and weakly-labelled datasets:

$$L = \frac{1}{n} \sum_{i=1}^{n} L_{\text{labelled}} + \frac{1}{m} \sum_{i=1}^{m} L_{\text{inferred}} \quad (1)$$

5. Repeat from step 2 using the current student model as the teacher model.

In our experiments, we compare this noisy self-training framework against the default 'non-noisy' self-training method, which simply skips step 3, meaning we do not apply any form of data augmentation.

3.3 Data Augmentation Methods

In each noisy self-training experiment we use nlpaug⁵ to apply one of the three following data augmentation methods for textual data:

Random Synonym Substitution Uses WordNet (Miller, 1995) to randomly replace tokens by one of its synonyms. For each sentence, 30% of its tokens will be replaced.

Random Word Swap Randomly swaps adjacent tokens in a sentence. For each sentence, 30% of its tokens are swapped.

Backtranslation First translates the original texts into a second language, then translates them back from the second language to the original language. We use the backtranslation model from nlpaug, which uses the two different transformer models from Ng et al. (2019) to translate the data from English to German, then from German back to English.

4 Experimental Setup

Firstly, we experiment with each dataset to estimate the hyperparameters for the base models, which is the first teacher models in the self-training loop. We use a batch size of 128, maximum sequence length of 128, learning rate of 0.00001, 15% of the training set as warm-up batches, weight decay of 0.001 and 20 training epochs. We apply a dropout rate of 10% for both the attention and classification layers. The model with highest validation F1-macro score⁶ obtained during training is loaded at the end of the last epoch. For the hyperparameters associated with the self-training method, we set the number of teacher-student iterations to 4 (including the first teacher model) and a confidence threshold filter of 80%, similarly to Xie et al. (2020a). Also, we experiment with five different pre-trained BERT models: DistilBERT, BERT-base-cased, BERT-large-cased, RoBERTabase and RoBERTa-large, aiming to investigate the impact of model size in performance gains associated with self-training.

From the above-listed configurations, we designed two main classification scenarios. The first scenario accounts for a regular self-training loop without data noise injection through augmentations, while the second scenario uses the noisy self-training approach, introducing data noise with one of the three augmentation methods described in Section 3.3.

Finally, we conduct a deeper analysis of each augmentation method. We use the first teacher model, trained exclusively with the human-labelled data of each dataset, to infer both the 'clean/without augmentation' and the 'noisy/augmented' versions of the unlabelled dataset and verify the following: (i) Does the augmentation method create new tokens that are not present in the vocabulary of the 'clean/without augmentation' unlabelled dataset? and (ii) Are the augmentations semantically invariant, meaning both the 'clean' and 'noisy' pairs of examples are assigned the same label?

5 Results

5.1 Default Fine-Tuning vs. Self-Training

Table 2 displays the mean and standard deviation F1-macro scores computed over three different random seed initializations for each experiment. Note

⁵https://github.com/makcedward/nlpaug

⁶Lowest training loss in the case of OLID, since no development set is provided.

DistilBERT 78.4 ± 0.1 79.2 ± 0.2 79.0 ± 0.3 79.0 ± 0.3 79.0 ± 0.3 79.0 BERT-base-cased 77.2 ± 0.3 78.7 ± 0.1 78.1 ± 0.1 78.3 ± 0.3 78.3 BERT-large-cased 79.2 ± 0.2 80.0 ± 0.3 79.4 ± 0.1 79.3 ± 0.3 79.3 RoBERTa-base 79.4 ± 0.7 80.1 ± 0.3 80.0 ± 0.4 $80.0 \pm 0.$						
DistilBERT 78.4 ± 0.1 79.2 ± 0.2 79.0 ± 0.3 78.3 ± 0.3 79.3 ± 0.3			OLID			
BERT-base-cased 77.2 ± 0.3 78.7 ± 0.1 78.1 ± 0.1 78.3 ± 0.3 78.3 BERT-large-cased 79.2 ± 0.2 80.0 ± 0.3 79.4 ± 0.1 79.3 ± 0.3 79.3 RoBERTa-base 79.4 ± 0.7 80.1 ± 0.3 80.0 ± 0.4 80.0 ± 0.4 80.0 RoBERTa-large 79.8 ± 0.3 80.4 ± 0.4 80.3 ± 0.4 80.7 ± 0.7 80.7 ConvAbuseArchitectureDFST $ST + BT$ $ST + SS$ $ST + SS$ DistilBERT 85.7 ± 0.5 86.8 ± 0.3 87.1 ± 0.3 87.2 ± 0.3 87.2 BERT-base-cased 86.8 ± 0.8 87.6 ± 0.1 87.2 ± 0.5 87.2 ± 0.5 87.2 BERT-large-cased 87.1 ± 0.6 87.9 ± 0.5 87.4 ± 0.2 87.9 ± 0.5 87.9 RoBERTa-base 84.5 ± 0.3 85.5 ± 0.4 85.3 ± 0.8 85.4 ± 0.5 85.4	Architecture	DF	ST	ST + BT	ST + SS	ST + WS
BERT-large-cased RoBERTa-base 79.2 ± 0.2 80.0 ± 0.3 79.4 ± 0.1 79.3 ± 0.3 79.3 RoBERTa-base RoBERTa-large 79.4 ± 0.7 80.1 ± 0.3 80.0 ± 0.4 80.0 ± 0.4 80.0 ± 0.4 80.0 ArchitectureDFSTST + BTST + SSST +DistilBERT 85.7 ± 0.5 86.8 ± 0.3 87.1 ± 0.3 87.2 ± 0.3 87.2 BERT-base-cased 86.8 ± 0.8 87.6 ± 0.1 87.2 ± 0.5 87.2 ± 0.5 87.2 ± 0.5 BERT-large-cased 87.1 ± 0.6 87.9 ± 0.5 87.4 ± 0.2 87.9 ± 0.5 87.9 ± 0.5 RoBERTa-base 84.5 ± 0.3 85.5 ± 0.4 85.3 ± 0.8 85.4 ± 0.5 85.4	DistilBERT	78.4 ± 0.1	79.2 ± 0.2	79.0 ± 0.3	79.0 ± 0.3	79.0 ± 0.3
RoBERTa-base 79.4 ± 0.7 80.1 ± 0.3 80.0 ± 0.4 80.0 ± 0.4 80.0 ± 0.4 RoBERTa-large 79.8 ± 0.3 80.4 ± 0.4 80.3 ± 0.4 80.7 ± 0.7 80.7 ConvAbuseArchitectureDFSTST + BTST + SSST +DistilBERT 85.7 ± 0.5 86.8 ± 0.3 87.1 ± 0.3 87.2 ± 0.3 87.2 BERT-base-cased 86.8 ± 0.8 87.6 ± 0.1 87.2 ± 0.5 87.2 ± 0.5 87.2 ± 0.5 BERT-large-cased 87.1 ± 0.6 87.9 ± 0.5 87.4 ± 0.2 87.9 ± 0.5 87.9 RoBERTa-base 84.5 ± 0.3 85.5 ± 0.4 85.3 ± 0.8 85.4 ± 0.5 85.4	BERT-base-cased	77.2 ± 0.3	78.7 ± 0.1	78.1 ± 0.1	78.3 ± 0.3	78.3 ± 0.3
RoBERTa-large 79.8 ± 0.3 80.4 ± 0.4 80.3 ± 0.4 80.7 ± 0.7 80.7 ConvAbuseArchitectureDFSTST + BTST + SSST +DistilBERT 85.7 ± 0.5 86.8 ± 0.3 87.1 ± 0.3 87.2 ± 0.3 87.2 ± 0.3 87.2 ± 0.5 BERT-base-cased 86.8 ± 0.8 87.6 ± 0.1 87.2 ± 0.5 BERT-large-cased 87.1 ± 0.6 87.9 ± 0.5 87.4 ± 0.2 87.9 ± 0.5 87.9 ± 0	BERT-large-cased	79.2 ± 0.2	80.0 ± 0.3	79.4 ± 0.1	79.3 ± 0.3	79.3 ± 0.3
ConvAbuseArchitectureDFSTST + BTST + SSST +DistilBERT 85.7 ± 0.5 86.8 ± 0.3 87.1 ± 0.3 87.2 ± 0.3 87.2 BERT-base-cased 86.8 ± 0.8 87.6 ± 0.1 87.2 ± 0.5 87.2 ± 0.5 87.2 ± 0.5 87.2 BERT-large-cased 87.1 ± 0.6 87.9 ± 0.5 87.4 ± 0.2 87.9 ± 0.5 87.9 ROBERTa-base 84.5 ± 0.3 85.5 ± 0.4 85.3 ± 0.8 85.4 ± 0.5 85.4	RoBERTa-base	79.4 ± 0.7	80.1 ± 0.3	80.0 ± 0.4	80.0 ± 0.4	80.0 ± 0.4
ArchitectureDFSTST + BTST + SSST +DistilBERT 85.7 ± 0.5 86.8 ± 0.3 87.1 ± 0.3 87.2 ± 0.3 87.2 ± 0.3 BERT-base-cased 86.8 ± 0.8 87.6 ± 0.1 87.2 ± 0.5 87.2 ± 0.5 87.2 ± 0.5 87.2 ± 0.5 BERT-large-cased 87.1 ± 0.6 87.9 ± 0.5 87.4 ± 0.2 87.9 ± 0.5 87.9 ± 0.5 ROBERTa-base 84.5 ± 0.3 85.5 ± 0.4 85.3 ± 0.8 85.4 ± 0.5 85.4	RoBERTa-large	79.8 ± 0.3	80.4 ± 0.4	80.3 ± 0.4	80.7 ± 0.7	80.7 ± 0.7
DistilBERT 85.7 ± 0.5 86.8 ± 0.3 87.1 ± 0.3 87.2 ± 0.3 87.2 BERT-base-cased 86.8 ± 0.8 87.6 ± 0.1 87.2 ± 0.5	ConvAbuse					
BERT-base-cased 86.8 ± 0.8 87.6 ± 0.1 87.2 ± 0.5 87.2 ± 0.5 87.2 ± 0.5 87.2 ± 0.5 BERT-large-cased 87.1 ± 0.6 87.9 ± 0.5 87.4 ± 0.2 87.9 ± 0.5 87.9 ± 0.5 ROBERTa-base 84.5 ± 0.3 85.5 ± 0.4 85.3 ± 0.8 85.4 ± 0.5 85.4 ± 0.5	Architecture	DF	ST	ST + BT	ST + SS	ST + WS
BERT-large-cased 87.1 \pm 0.6 87.9 \pm 0.5 87.4 \pm 0.2 87.9 \pm 0.5 87.9 RoBERTa-base 84.5 \pm 0.3 85.5 \pm 0.4 85.3 \pm 0.8 85.4 \pm 0.5 85.4	DistilBERT	85.7 ± 0.5	86.8 ± 0.3	87.1 ± 0.3	87.2 ± 0.3	87.2 ± 0.3
RoBERTa-base 84.5 ± 0.3 85.5 ± 0.4 85.3 ± 0.8 85.4 ± 0.5 85.4	BERT-base-cased	86.8 ± 0.8	87.6 ± 0.1	87.2 ± 0.5	87.2 ± 0.5	87.2 ± 0.5
	BERT-large-cased	87.1 ± 0.6	87.9 ± 0.5	87.4 ± 0.2	87.9 ± 0.5	87.9 ± 0.5
RoBERTa-large 86.0 ± 0.1 86.2 ± 0.3 86.6 ± 0.3 86.9 ± 0.1 86.8	RoBERTa-base	84.5 ± 0.3	85.5 ± 0.4	85.3 ± 0.8	85.4 ± 0.5	85.4 ± 0.5
	RoBERTa-large	86.0 ± 0.1	86.2 ± 0.3	86.6 ± 0.3	86.9 ± 0.1	86.8 ± 0.1

Table 2: Mean ± 1 std F1-Macro scores obtained over three random seed initializations.

DF=Default Fine-Tuning, ST=Self-Training, BT=Backtranslation, SS=Synonym Substitution, WS=Word Swap

that self-training, regardless of whether coupled with data augmentation methods or not, improves over default fine-tuning for every model architecture, increasing the F1-macro score from +0.7% up to +1.5% on OLID and +0.8% up to +1.5% on ConvAbuse depending on the pre-trained model architecture.

Also, we highlight how self-training can make smaller models, which require fewer resources to maintain in practical applications, achieving the same performance as larger and more costly models that are trained with default fine-tuning. Self-training on a DistilBERT (66M parameters) outperforms a BERT-large-cased (340M parameters) with default fine-tuning on both OLID and ConvAbuse. On OLID, a RoBERTa-base architecture (125M parameters) with self-training outperforms a RoBERTa-large (354M parameters) architecture with default fine-tuning, although this does not hold true for ConvAbuse.

Furthermore, we point out that OLID and ConvAbuse's data come from different sources, the first being Twitter, and the second one representing conversations between humans and chatbots, thus their structure differs significantly. Since our unlabelled dataset is composed of Twitter data, it would be fair to assume that the benefits of self-training in our experiments would be more prominent for the OLID dataset, but our results do not show this, since models trained with ConvAbuse benefited from self-training with our Twitter-originated unlabelled dataset just as much as models trained with

OLID.

5.2 Default Self-Training vs. Noisy Self-Training

After verifying that self-training is beneficial to both datasets on all model architectures, we compare default self-training with noisy self-training, and the impacts of adding data noise in the form of data augmentations. We find that introducing data augmentations to the self-training pipeline increases performance against default self-training only for RoBERTa-large on both OLID and ConvAbuse, with DistilBERT also showing improvements for ConvAbuse, but not for OLID. On all other architectures, for both datasets, default self-training without data augmentations achieves the highest scores.

In our results for offensive/hate speech classification, backtranslation does not achieve the highest score in any setup, while synonym substitution and word swap tie for highest score in three scenarios: ConvAbuse with DistilBERT, ConvAbuse with BERT-large-cased, and OLID with RoBERTa-large. Synonym substitution outperforms all the remaining methods on ConvAbuse with RoBERTa-large.

An important remark is that our results diverge from He et al. (2019), which finds that state-of-theart data augmentation methods such as backtranslation outperform simpler methods on self-training for machine translation and text summarization. However, our results align with Marivate and Sefara (2020), although their work is not focused on self-training, but instead on how different data augmentation techniques impact their models on three datasets from different domains. They report backtranslation as their worst augmentation method on a hate speech dataset, decreasing accuracy by -2.3%. Our findings bridge this gap and reveal that backtranslation has significant limitations in the domain of offensive/hate speech detection, even when used in a noisy self-training approach.

5.3 Data Augmentation Analysis

Our first data augmentation analysis is to understand if the augmented text introduces new unseen tokens to the vocabulary of the 'clean' unlabelled set when both are combined. We find a vocabulary size increase of 39.5%, 9.0% and 4.7% averaging across all different pre-trained architectures for backtranslation, synonym substitution and word swap⁷ respectively. This indicates that backtranslation is heavily superior in terms of introducing new unseen tokens, but this is not correlated with performance increase, as backtranslation appears as the worst augmentation method for noisy self-training in our classification experiments.

Next, in order to verify the performance of the data augmentation methods in generating semantically invariant examples, we use the base models trained exclusively with the human-labelled data from each dataset, on each pre-trained architecture, and use them to perform inference on both the 'clean' and the noisy/augmented unlabelled set. We then compare both predictions and analyse how augmentations may shift the underlying target class. We will refer to **positive shift** when a non-offensive example is classified as offensive after being augmented, and **negative shift** when an offensive example is classified as non-offensive after being augmented.

Table 3 presents the total class shift percentage for each augmentation method, averaging across both datasets and all model architectures, of which we further divide into positive and negative label shift percentages. Notice that backtranslation is the method that produces the highest amount of label shifting at 23.8%, of which 54.7% are negative shifts, which is a 6.6% increase over synonym substitution and a 4.8% increase word swap.

It is fair to assume that not all of the class shifting occurs from the augmentation changing the seman-

Augmentation	Total Shift	Positive Shift	Negative Shift
BT	23.8%	46.7%	54.7%
SS	23.5%	48.7%	51.3%
WS	23.3%	47.8%	52.2%

Table 3: Average target class shift percentage on the weakly-labelled set. BT=Backtranslation, SS=Synonym Substitution, WS=Word Swap

tic that defines if an example is either offensive or not-offensive. In most cases, class shifting may occur because of small perturbations that are semantically invariant, meaning both the 'clean' and the augmented text's true underlying classes are still the same, even if the classifier predicted them as different classes. In these cases, when we set the label of the augmented text to be the same as the one obtained when inferring the 'clean' version of the text, as presented in section 3.2, we are reinforcing the model to be more robust against these small perturbations, which is one of the main benefits of noisy self-training. However, when augmentation methods create semantically different versions of the original texts, replicating the inferred label from the original text to the augmented text results in the addition of incorrect ground-truth labels to the train set, which may degrade performance.

Currently, to the best of our knowledge, there is no dataset annotated for offense/hate speech before and after applying data augmentation, which would enable a more accurate estimation of semantic variations produced by them. In tables 4 and 5 we show two examples for each augmentation method that suffered from positive shift (not-offensive to offensive) and negative shift (offensive to not-offensive), respectively.

An example of a recurrent theme among various target shifted examples is the substitution of the keywords 'fuck' with 'damn' or 'hell', indicating that despite these keywords being semantically similar, they are not always interchangeable with respect to the target class, and the mere replacement of one for another is enough to shift the target class. This could be expected, as offense detection is highly impacted by the mere presence or absence of offensive keywords.

6 Conclusion

In this work, we analysed the impact of selftraining on offensive and hate speech classification tasks using five different pre-trained BERT models

⁷Word swap is unintuitively capable of creating new tokens depending on how a sentence is split into tokens and then merged back after swapping the tokens.

Text	Augmented Text	Method
I HATE ALL OF YOU	ALL I HATE OF YOU	WS
Maybe I dont respect all women	Maybe I respect dont women all	WS
Bitches and sports	Females and Sport	BT
Wooooow what the fuck	Wooooow, what the hell?	BT
Bitch you better be joking	Gripe you good be joking	SS
The NYT has been showing its whole ass []	The NYT has follow showing its whole butt []	SS

Table 4: Examples of Offensive to Not-Offensive semantic shift created by data augmentation. BT=Backtranslation, SS=Synonym Substitution, WS=Word Swap

Text	Augmented Text	Method
Is that Fat Albert	That Fat is Albert	WS
Man that is terrible	That man is terrible	WS
damn white people oppressing the blacks	fucking white people who oppress the blacks	BT
That damn staircase be beating my ass []	That fucking staircase will bang my ass []	BT
i will not get over this	i will not fuck off ended this	SS
Send me the link and Ill love you forever	Send pine tree state the link and Ill fuck you forever	SS

Table 5: Examples of Not-Offensive to Offensive class shift created by data augmentation. BT=Backtranslation, SS=Synonym Substitution, WS=Word Swap

of varying sizes and two different datasets. We also experimented with noisy self-training using three different data augmentation techniques for textual data. We found that self-training improves classification performance for all model architectures on both datasets, with an increase in F1-Macro of up to +1.5%. However, our experiments comparing default self-training versus noisy self-training showed that noisy self-training does not improve performance, despite its success in other domains. Finally, we investigated the three data augmentation methods and showed that the domain of offensive/hate speech classification is highly sensitive to semantic variances produced by them, and we discussed future research ideas to mitigate these problems.

7 Future Work

We understand that some of the semantic variations discussed in this work could be mitigated by data augmentation methods that both preserve existing offensive keywords, and do not introduce new offensive keywords randomly, as these are often conditional to the underlying ground-truth class. For some languages, most of these keywords are extensively documented⁸, thus they can be known a priori by these methods, and be treated differently, such as only substituting an offensive keyword by

Acknowledgments

We thank Olesya Razuvayevskaya and Freddy Heppell for their valuable feedback. This research has been funded by "SoBigData++: European Integrated Infrastructure for Social Mining and Big Data Analytics" (EU H2020, Grant Agreement n.871042 (http://www.sobigdata.eu)).

another offensive keyword, or not allowing a nonoffensive keyword to be substituted by an offensive keyword. This custom approach can theoretically help mitigate semantic variations in this domain, but offensive/hateful comments can still be made without making use of a single offensive/hateful keyword. In these more subtle cases, a system would have to detect the offensive/hateful context without relying solely on keywords, and modify the example while still maintaining this context. We see potential benefits of using recent instructiontuned large language models (Ouyang et al., 2022) as specialised data augmentation methods that are task-specific, and can be able to preserve the semantics associated with the task when modifying a given text. In this scenario, an instruction prompt can be designed to inform the system of the context of the task, and make it aware that this semantic must be preserved when modifying the given text. In the future, we aim towards extending this work with the above-mentioned research ideas.

⁸https://hatebase.org/

References

- Jean-Baptiste Alayrac, Jonathan Uesato, Po-Sen Huang, Alhussein Fawzi, Robert Stanforth, and Pushmeet Kohli. 2019. Are labels required for improving adversarial robustness? In Advances in Neural Information Processing Systems, volume 32. Curran Associates, Inc.
- Safa Alsafari and Samira Sadaoui. 2021. Semisupervised self-training of hate and offensive speech from social media. *Applied Artificial Intelligence*, 35(15):1621–1645.
- Massih-Reza Amini, Vasilii Feofanov, Loic Pauletto, Emilie Devijver, and Yury Maximov. 2022. Self-training: A survey. *arXiv preprint arXiv:2202.12040*.
- Markus Bayer, Marc-André Kaufhold, and Christian Reuter. 2022. A survey on data augmentation for text classification. *ACM Computing Surveys*, 55(7):1–39.
- Avrim Blum and Tom Mitchell. 1998. Combining labeled and unlabeled data with co-training. In *Proceedings of the Eleventh Annual Conference on Computational Learning Theory*, COLT' 98, page 92–100, New York, NY, USA. Association for Computing Machinery.
- Yair Carmon, Aditi Raghunathan, Ludwig Schmidt, John C Duchi, and Percy S Liang. 2019. Unlabeled data improves adversarial robustness. In *Advances in Neural Information Processing Systems*, volume 32. Curran Associates, Inc.
- Amanda Cercas Curry, Gavin Abercrombie, and Verena Rieser. 2021. ConvAbuse: Data, analysis, and benchmarks for nuanced abuse detection in conversational AI. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 7388–7403, Online and Punta Cana, Dominican Republic. Association for Computational Linguistics.
- X. Chen, Y. Yuan, G. Zeng, and J. Wang. 2021. Semisupervised semantic segmentation with cross pseudo supervision. In 2021 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), pages 2613–2622, Los Alamitos, CA, USA. IEEE Computer Society.
- E. D. Cubuk, B. Zoph, J. Shlens, and Q. V. Le. 2020. Randaugment: Practical automated data augmentation with a reduced search space. In 2020 IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW), pages 3008–3017, Los Alamitos, CA, USA. IEEE Computer Society.
- Thomas Davidson, Dana Warmsley, Michael Macy, and Ingmar Weber. 2017. Automated hate speech detection and the problem of offensive language. *Proceedings of the International AAAI Conference on Web and Social Media*, 11(1):512–515.
- J. Deng, W. Dong, R. Socher, L. Li, Kai Li, and Li Fei-Fei. 2009. Imagenet: A large-scale hierarchical image database. In 2009 IEEE Computer Society Conference on Computer Vision and Pattern Recognition

- Workshops (CVPR Workshops), pages 248–255, Los Alamitos, CA, USA. IEEE Computer Society.
- Sergey Edunov, Myle Ott, Michael Auli, and David Grangier. 2018. Understanding back-translation at scale. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 489–500, Brussels, Belgium. Association for Computational Linguistics.
- Paula Fortuna and Sérgio Nunes. 2018. A survey on automatic detection of hate speech in text. ACM Comput. Surv., 51(4).
- Alec Go, Richa Bhayani, and Lei Huang. 2009. Twitter sentiment classification using distant supervision. *CS224N project report, Stanford*, 1(12):2009.
- Junxian He, Jiatao Gu, Jiajun Shen, and Marc'Aurelio Ranzato. 2019. Revisiting Self-Training for Neural Sequence Generation. arXiv e-prints, page arXiv:1909.13788.
- Mai Ibrahim, Marwan Torki, and Nagwa El-Makky. 2018. Imbalanced toxic comments classification using data augmentation and deep learning. In 2018 17th IEEE International Conference on Machine Learning and Applications (ICMLA), pages 875–878.
- Giannis Karamanolakis, Subhabrata (Subho) Mukherjee, Guoqing Zheng, and Ahmed H. Awadallah. 2021. Self-training with weak supervision. In *NAACL* 2021. NAACL 2021.
- Samuli Laine and Timo Aila. 2017. Temporal ensembling for semi-supervised learning. In 5th International Conference on Learning Representations, ICLR 2017, Toulon, France, April 24-26, 2017, Conference Track Proceedings. OpenReview.net.
- Elisa Leonardelli, Stefano Menini, and Sara Tonelli. 2020. Dh-fbk@ haspeede2: Italian hate speech detection via self-training and oversampling. In *Proceedings of the Seventh Evaluation Campaign of Natural Language Processing and Speech Tools for Italian. Final Workshop (EVALITA 2020)*, volume 2765.
- Shuo Liu, Adria Mallol-Ragolta, Emilia Parada-Cabaleiro, Kun Qian, Xin Jing, Alexander Kathan, Bin Hu, and Björn W. Schuller. 2022. Audio self-supervised learning: A survey. *Patterns*, 3(12):100616.
- Shayne Longpre, Yu Wang, and Chris DuBois. 2020. How effective is task-agnostic data augmentation for pretrained transformers? In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 4401–4411, Online. Association for Computational Linguistics.
- Vukosi Marivate and Tshephisho Sefara. 2020. Improving short text classification through global augmentation methods. In *Machine Learning and Knowledge Extraction:* 4th IFIP TC 5, TC 12, WG 8.4, WG 8.9, WG 12.9 International Cross-Domain Conference, CD-MAKE 2020, Dublin, Ireland, August 25–28, 2020, Proceedings 4, pages 385–399. Springer.

- George A Miller. 1995. Wordnet: a lexical database for english. *Communications of the ACM*, 38(11):39–41.
- Takeru Miyato, Shin-ichi Maeda, Masanori Koyama, and Shin Ishii. 2018. Virtual adversarial training: a regularization method for supervised and semi-supervised learning. *IEEE transactions on pattern analysis and machine intelligence*, 41(8):1979–1993.
- Mainack Mondal, Leandro Araújo Silva, and Fabrício Benevenuto. 2017. A measurement study of hate speech in social media. In *Proceedings of the 28th ACM Conference on Hypertext and Social Media*, HT '17, page 85–94, New York, NY, USA. Association for Computing Machinery.
- Anna Mosolova, Vadim Fomin, and Ivan Bondarenko. 2018. Text augmentation for neural networks. *AIST* (*Supplement*), 2268:104–109.
- Amir Najafi, Shin-ichi Maeda, Masanori Koyama, and Takeru Miyato. 2019. Robustness to adversarial perturbations in learning from incomplete data. In *Advances in Neural Information Processing Systems*, volume 32. Curran Associates, Inc.
- Nathan Ng, Kyra Yee, Alexei Baevski, Myle Ott, Michael Auli, and Sergey Edunov. 2019. Facebook FAIR's WMT19 news translation task submission. In *Proceedings of the Fourth Conference on Machine Translation (Volume 2: Shared Task Papers, Day 1)*, pages 314–319, Florence, Italy. Association for Computational Linguistics.
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. 2022. Training language models to follow instructions with human feedback. Advances in Neural Information Processing Systems, 35:27730–27744.
- Bao-Tran Pham-Hong and Setu Chokshi. 2020. PGSG at SemEval-2020 task 12: BERT-LSTM with tweets' pretrained model and noisy student training method. In *Proceedings of the Fourteenth Workshop on Semantic Evaluation*, pages 2111–2116, Barcelona (online). International Committee for Computational Linguistics.
- Antti Rasmus, Harri Valpola, Mikko Honkala, Mathias Berglund, and Tapani Raiko. 2015. Semi-supervised learning with ladder networks. In *Proceedings of the 28th International Conference on Neural Information Processing Systems Volume 2*, NIPS'15, page 3546–3554, Cambridge, MA, USA. MIT Press.
- Caitlin Richardson, Sandeep Shah, and Xiaohong Yuan. 2022. Semi-supervised machine learning for analyzing covid-19 related twitter data for asian hate speech. In 2022 21st IEEE International Conference on Machine Learning and Applications (ICMLA), pages 1643–1648. IEEE.
- Georgios Rizos, Konstantin Hemker, and Björn Schuller. 2019. Augment to prevent: Short-text data augmentation in deep learning for hate-speech classification.

- In Proceedings of the 28th ACM International Conference on Information and Knowledge Management, CIKM '19, page 991–1000, New York, NY, USA. Association for Computing Machinery.
- Raquel Bento Santos, Bernardo Cunha Matos, Paula Carvalho, Fernando Batista, and Ricardo Ribeiro. 2022. Semi-Supervised Annotation of Portuguese Hate Speech Across Social Media Domains. In 11th Symposium on Languages, Applications and Technologies (SLATE 2022), volume 104 of Open Access Series in Informatics (OASIcs), pages 11:1–11:14, Dagstuhl, Germany. Schloss Dagstuhl Leibniz-Zentrum für Informatik.
- Madeline C. Schiappa, Yogesh S. Rawat, and Mubarak Shah. 2022. Self-supervised learning for videos: A survey. *ACM Comput. Surv.* Just Accepted.
- Zeerak Waseem. 2016. Are you a racist or am I seeing things? annotator influence on hate speech detection on Twitter. In *Proceedings of the First Workshop on NLP and Computational Social Science*, pages 138–142, Austin, Texas. Association for Computational Linguistics.
- Zeerak Waseem and Dirk Hovy. 2016. Hateful symbols or hateful people? predictive features for hate speech detection on Twitter. In *Proceedings of the NAACL Student Research Workshop*, pages 88–93, San Diego, California. Association for Computational Linguistics.
- Qizhe Xie, Zihang Dai, Eduard Hovy, Minh-Thang Luong, and Quoc V. Le. 2020a. Unsupervised data augmentation for consistency training. In *Proceedings of the 34th International Conference on Neural Information Processing Systems*, NIPS'20, Red Hook, NY, USA. Curran Associates Inc.
- Qizhe Xie, Minh-Thang Luong, Eduard Hovy, and Quoc V Le. 2020b. Self-training with noisy student improves imagenet classification. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 10687–10698.
- Marcos Zampieri, Shervin Malmasi, Preslav Nakov, Sara Rosenthal, Noura Farra, and Ritesh Kumar. 2019. Predicting the type and target of offensive posts in social media. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 1415–1420, Minneapolis, Minnesota. Association for Computational Linguistics.
- Hongyi Zhang, Moustapha Cisse, Yann N. Dauphin, and David Lopez-Paz. 2018. mixup: Beyond empirical risk minimization. In *International Conference on Learning Representations*.
- Xiang Zhang, Junbo Zhao, and Yann LeCun. 2015. Character-level convolutional networks for text classification. In *Advances in Neural Information Processing Systems*, volume 28. Curran Associates, Inc.